



CUDA编程基础 – 利用CUDA加速图像处理

NVIDIA企业级开发者社区 何琨

AGENDA

CUDA并行计算基础

- GPU硬件平台
- CUDA 安装
- CUDA线程层次
- CUDA线程索引
- CUDA线程分配
- 图像处理
- 编程实例：Sobel边缘检测

What is CUDA?

- CUDA
 - Compute Unified Device Architecture
- CUDA C/C++
 - 基于C/C++的编程方法
 - 支持异构编程的扩展方法
 - 简单明了的APIs，能够轻松的管理存储系统
- CUDA支持的编程语言：
 - C/C++/Python/Fortran/Java/.....

NVIDIA 开发者

15

500k


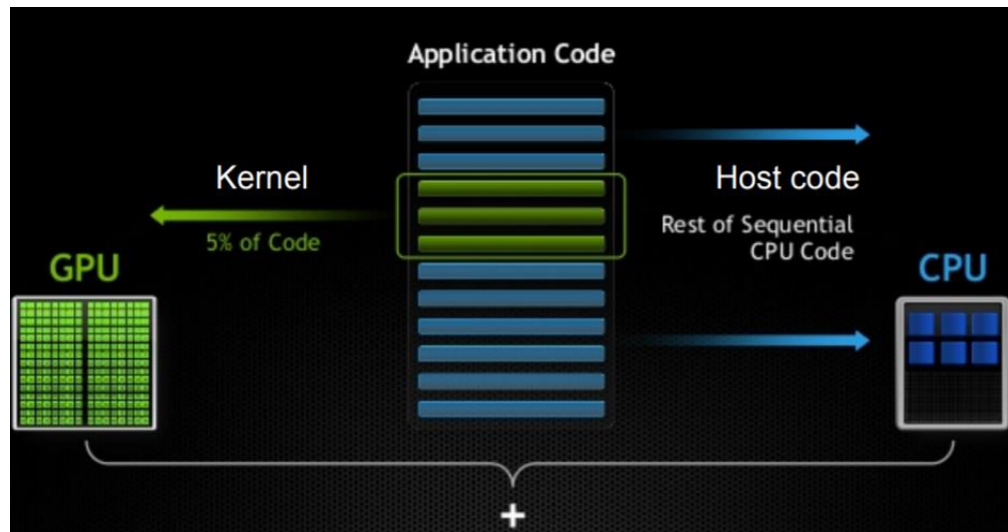


**MILLION
DEVELOPERS**

异构计算

- 术语:

- *Host* CPU和内存(host memory)
- *Device* GPU和显存(device memory)




Host



Device


异构计算

高性能计算大会ISCTOP 500



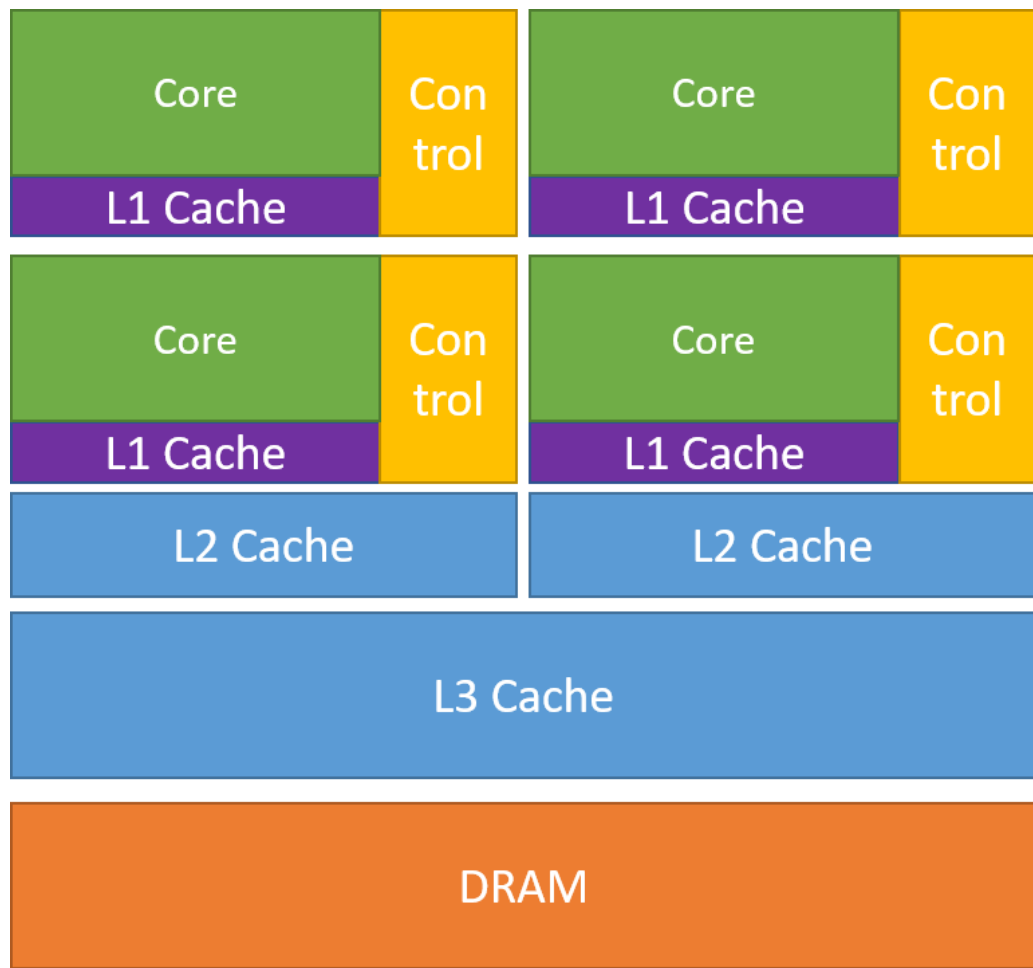
■ GPU ■ none GPU

高性能计算大会ISC TOP10




■ GPU ■ none GPU

芯片结构

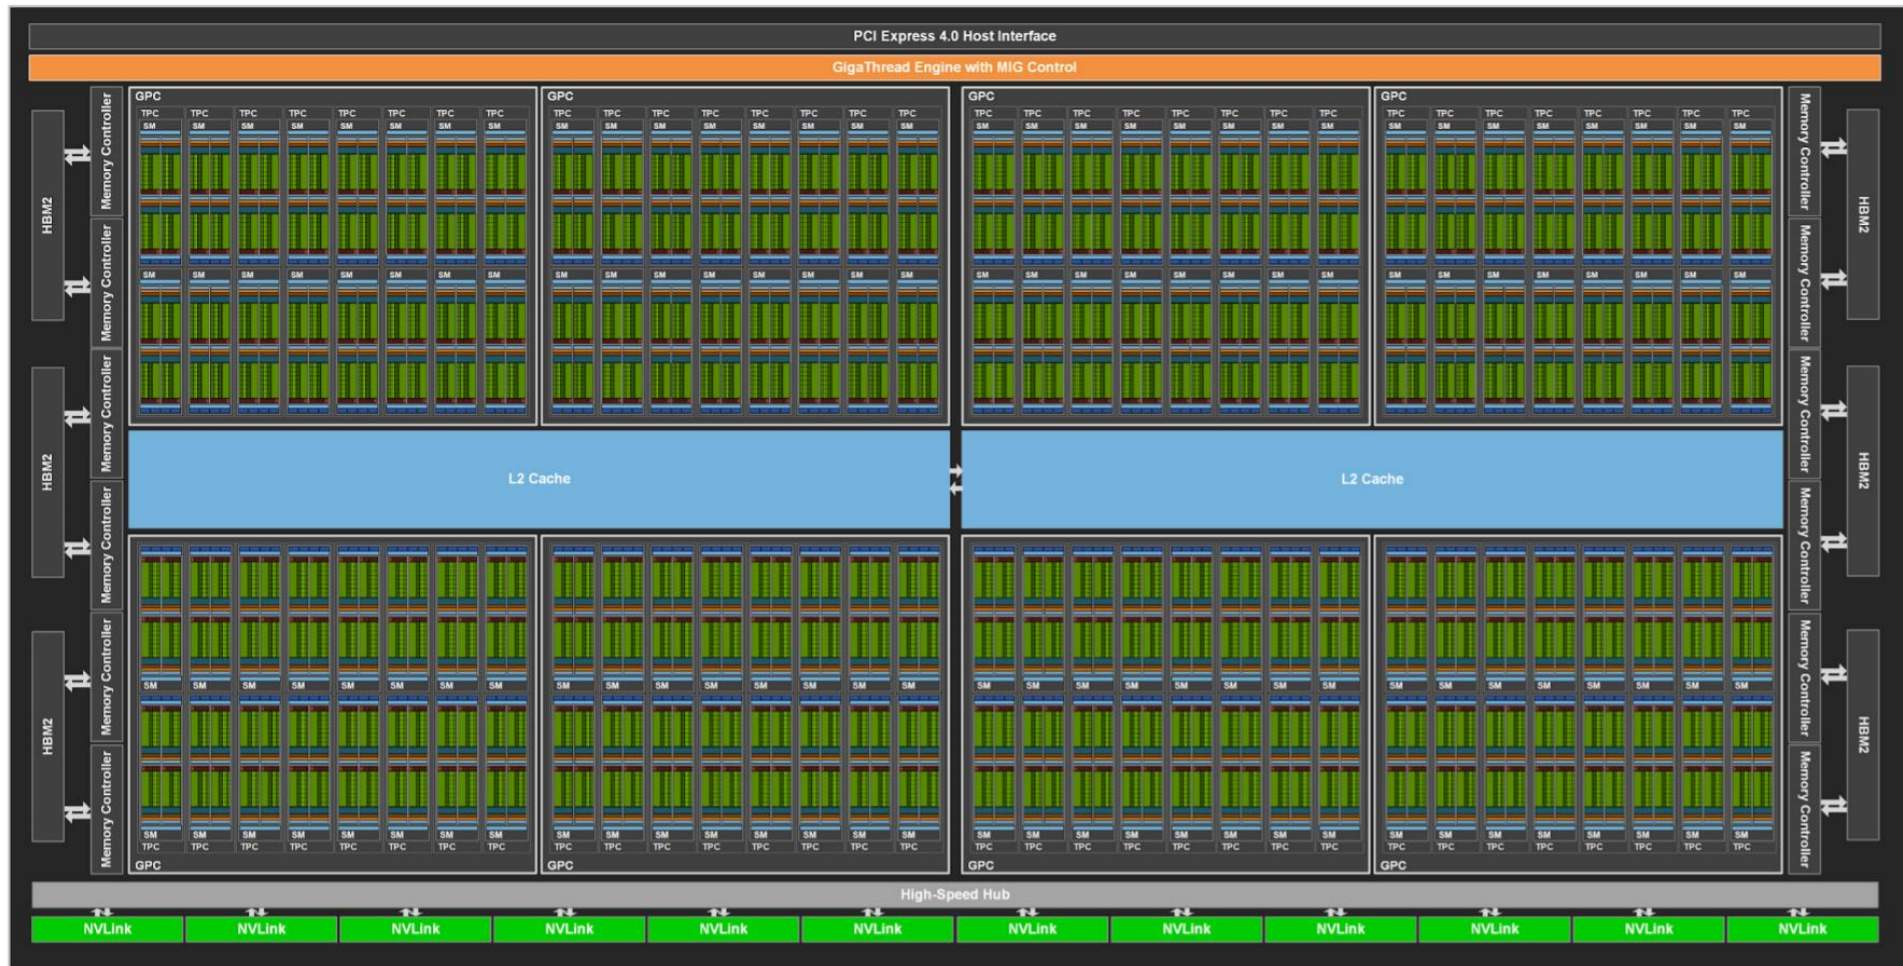


CPU



GPU

GPU结构---GA100



GPU结构---GA100



- 8 GPCs, 8 TPCs/GPC, 2 SMs/TPC, 16 SMs/GPC, 128 SMs per full GPU
- 64 FP32 CUDA Cores/SM, 8192 FP32 CUDA Cores per full GPU
- 4 third-generation Tensor Cores/SM, 512 third-generation Tensor Cores per full GPU
- 6 HBM2 stacks, 12 512-bit memory controllers


CUDA安装

- 适用设备:
 - 所有包含NVIDIA GPU的服务器，工作站，个人电脑，嵌入式设备等电子设备
- 软件安装:
 - Windows: <https://docs.nvidia.com/cuda/cuda-installation-guide-microsoft-windows/index.html>
只需安装一个.exe的可执行程序
 - Linux: <https://docs.nvidia.com/cuda/cuda-installation-guide-linux/index.html>
按照上面的教程，需要6 / 7 个步骤即可
 - Jetson: <https://developer.nvidia.com/embedded/jetpack>
直接利用NVIDIA SDK Manager 或者 SD image进行刷机即可

CUDA线程层次

```
HelloFromGPU <<<?, ?>>>();
```

- ❖ **Thread**: sequential execution unit
 - 所有线程执行相同的核函数
 - 并行执行
- ❖ **Thread Block**: a group of threads
 - 执行在一个Streaming Multiprocessor (SM)
 - 同一个Block中的线程可以协作
- ❖ **Thread Grid**: a collection of thread blocks
 - 一个Grid当中的Block可以在多个SM中执行



线程层次

❖ 执行设置:

`dim3 grid(3,2,1), block(5,3,1)`

❖ Built-in variables:

- `threadIdx.[x y z]`

是执行当前kernel函数的线程在block中的索引值

- `blockIdx.[x y z]`

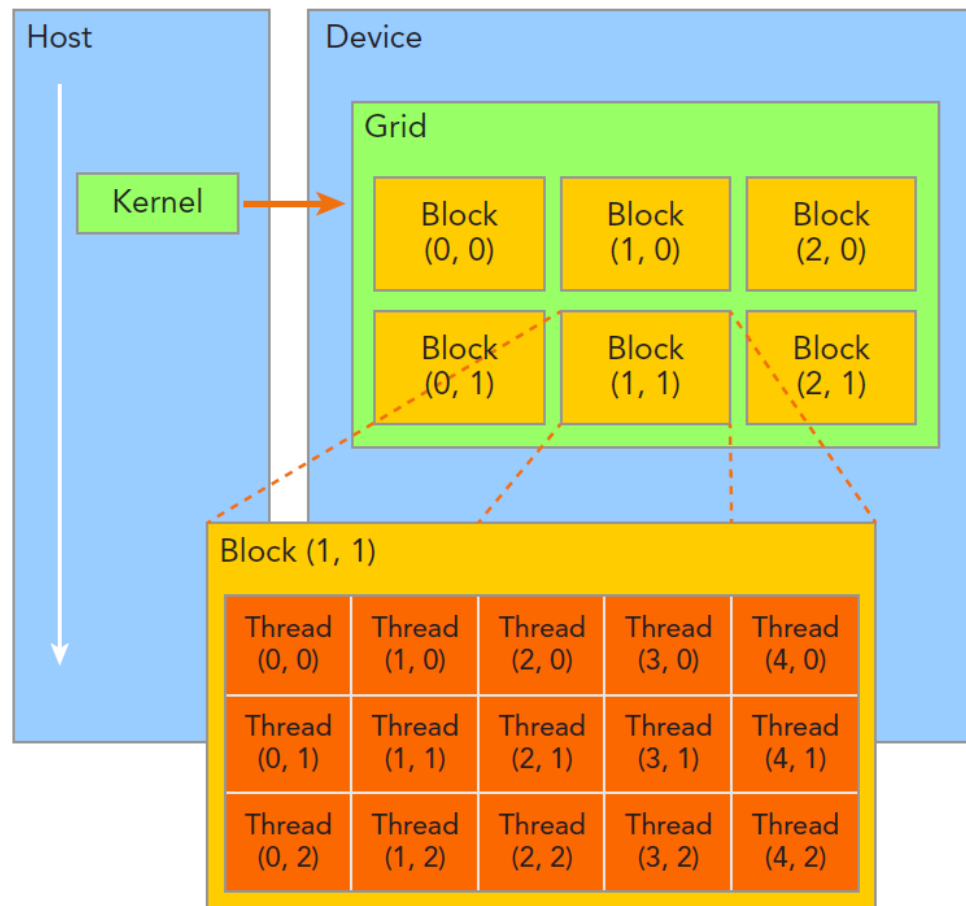
是指执行当前kernel函数的线程所在block，在grid中的索引值

- `blockDim.[x y z]`

表示一个grid中包含多少个block


- `gridDim.[x y z]`

表示一个block中包含多少个线程



CUDA的线程索引

- 如何确定线程执行地数据



```
int index = threadIdx.x + blockIdx.x * blockDim.x;  
          =          5      +          2      * 8;  
          = 21;
```


线程层次

- 我们写的程序:

```
__global__ void add( int *a, int *b, int *c ) {  
    c[threadIdx.x] = a[threadIdx.x] + b[threadIdx.x];  
}
```

```
add<<<1,4>>>( a, b, c);
```

- 实际上在设备上运行的样子:

Thread 0

```
c[0] = a[0] + b[0];
```

Thread 1

```
c[1] = a[1] + b[1];
```

Thread 2

```
c[2] = a[2] + b[2];
```

Thread 3

```
c[3] = a[3] + b[3];
```

线程层次

Software

GPU




Thread




CUDA Core

Threads are executed by cuda core



SM

Thread blocks are executed on SM



Grid

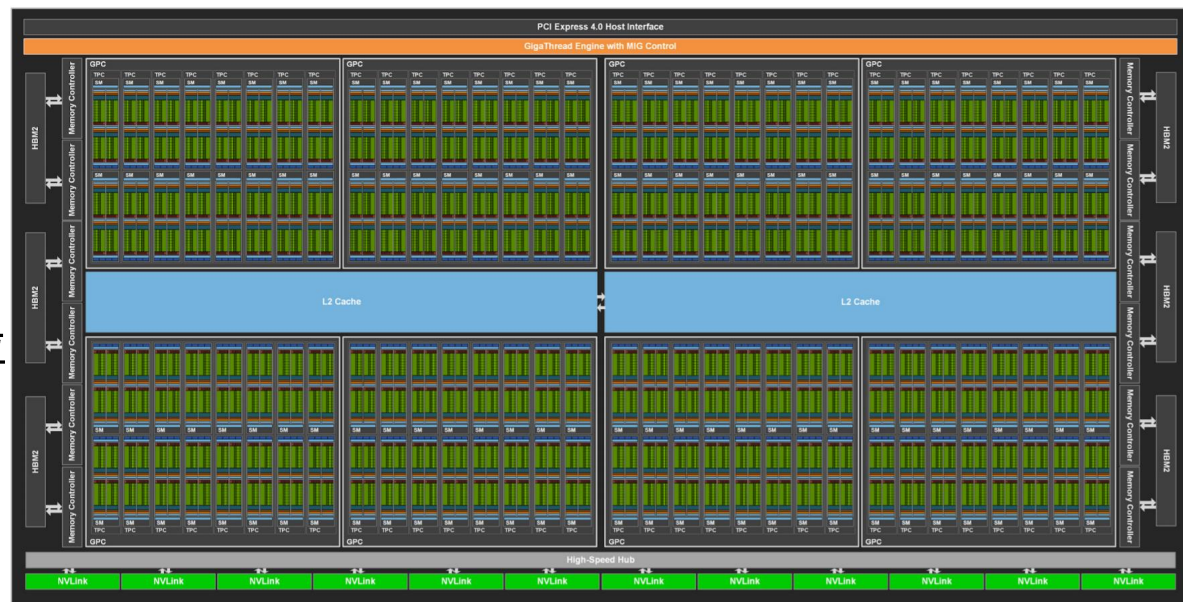


Device

A kernel is launched as a grid of thread blocks

线程层次

- 硬件调度：
 - Grid: GPU(GPC)级别的调度单位
 - Block(CTA): SM级别的调度单位
 - Threads/Warp: CUDA core级别的调度单位
- 资源和通信：
 - Grid: 共享同样的kernel 和 Context
 - Block(CTA): 同一个SM(Streaming Multiprocessor), 同一个SM(Shared Memory)
 - Threads/Warp: 允许同一个warp中的thread读取其他thread的值



线程层次

- 硬件调度：
 - Grid: GPU(GPC)级别的调度单位
 - Block(CTA): SM级别的调度单位
 - Threads/Warp: CUDA core级别的调度单位
- 资源和通信：
 - Grid: 共享同样的kernel 和 Context
 - Block(CTA): 同一个SM(Streaming Multiprocessor), 同一个SM(Shared Memory)
 - Threads/Warp: 允许同一个warp中的thread读取其他thread的值



CUDA的线程索引

```
__global__ void add(const double *x, const double  
*y, double *z)  
{  
    const int n = blockDim.x * blockIdx.x + threadIdx.x;  
    z[n] = x[n] + y[n];  
}
```

每个线程都执行相同的命令

CUDA PROGRAMMING BY EXAMPLE


Case: Vector Add

❖ Parallelizable problem:

- $c = a + b$
- a, b, c are vectors of length N

❖ CPU implementation:

```
void main(){  
    int size = N * sizeof(int);  
    int *a, *b, *c;  
    a = (int *)malloc(size);  
    b = (int *)malloc(size);  
    c = (int *)malloc(size);  
    memset(c, 0, size);  
    init_rand_f(a, N);  
    init_rand_f(b, N);  
  
    vecAdd(N, a, b, c);  
}
```



```
void vecAdd (int n, int *a,  
            int *b, int *c)  
{  
    for(int i = 0; i < n; i++)  
    {  
        c[i] = a[i] + b[i];  
    }  
}
```

GPU CODE WORKFLOW

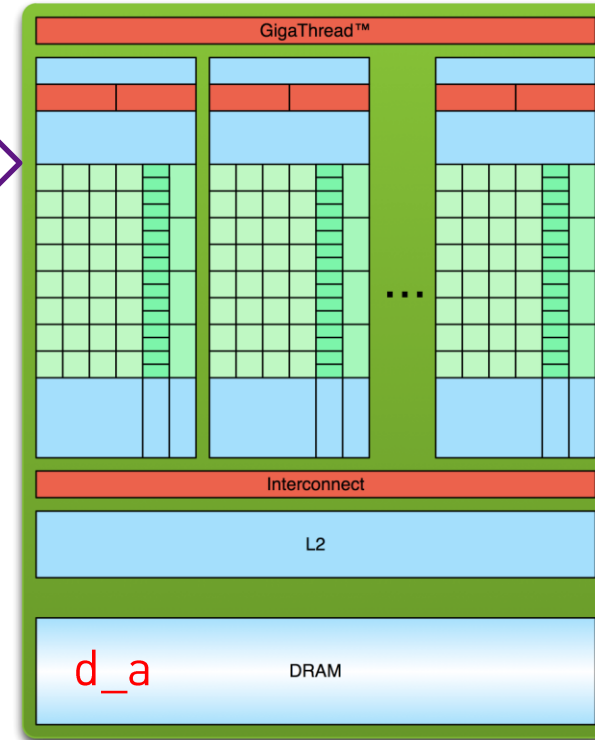
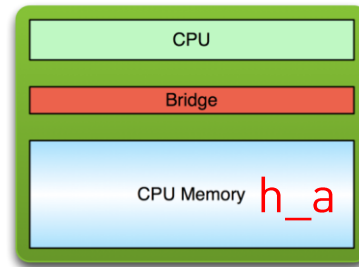
Allocate GPU Memories

```
int main(void) {  
    size_t size = N * sizeof(int);  
    int *h_a, *h_b; int *d_a, *d_b, *d_c;  
    h_a = (int *)malloc(size);  
    h_b = (int *)malloc(size);  
    ...
```

```
    cudaMalloc((void **)&d_a, size);  
    cudaMalloc((void **)&d_b, size);  
    cudaMalloc((void **)&d_c, size);
```

```
    cudaMemcpy(d_a, h_a, size, cudaMemcpyHostToDevice);  
    cudaMemcpy(d_b, h_b, size, cudaMemcpyHostToDevice);  
    vectorAdd<<<grid, block>>>(d_a, d_b, d_c, N);  
    cudaMemcpy(h_c, d_c, size, cudaMemcpyDeviceToHost);
```

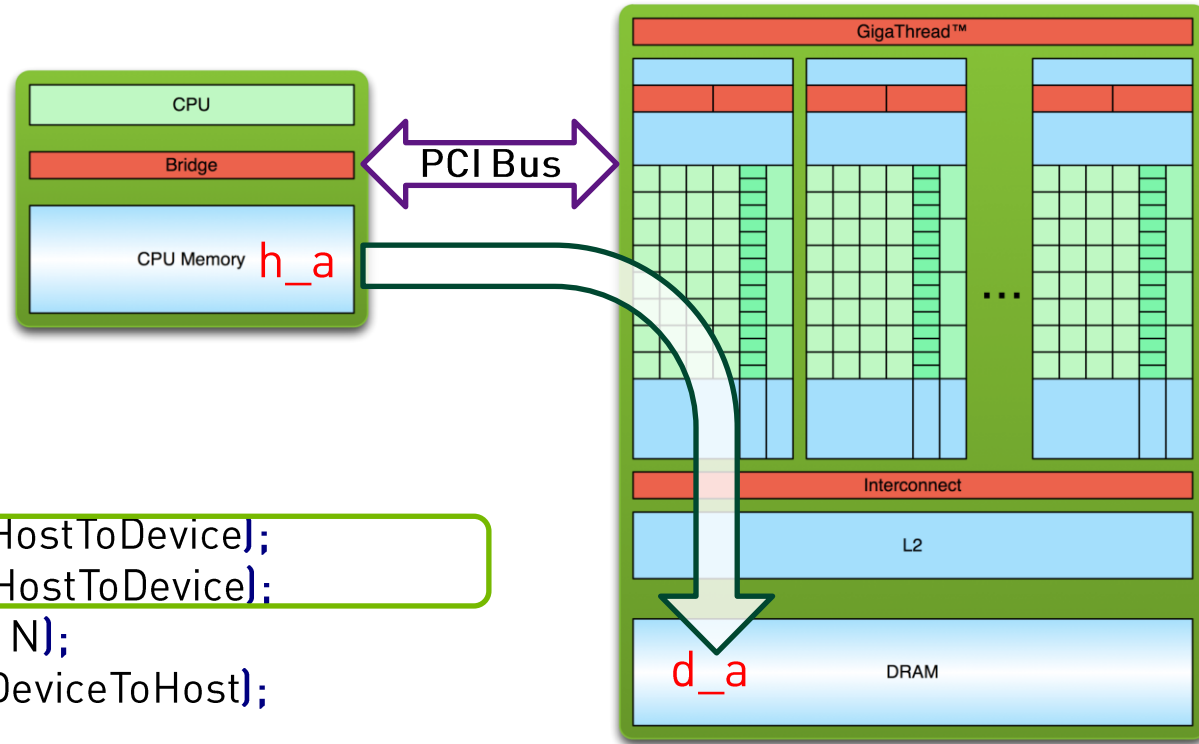
```
    cudaFree(d_a); cudaFree(d_b); cudaFree(d_c);  
    free(h_a); free(h_b);  
    return 0;}
```



GPU CODE WORKFLOW

Copy data from CPU to GPU

```
int main(void) {  
    size_t size = N * sizeof(int);  
    int *h_a, *h_b; int *d_a, *d_b, *d_c;  
    h_a = (int *)malloc(size);  
    h_b = (int *)malloc(size);  
    ...  
    cudaMalloc((void **)&d_a, size);  
    cudaMalloc((void **)&d_b, size);  
    cudaMalloc((void **)&d_c, size);  
  
    cudaMemcpy(d_a, h_a, size, cudaMemcpyHostToDevice);  
    cudaMemcpy(d_b, h_b, size, cudaMemcpyHostToDevice);  
    vectorAdd<<<grid, block>>>(d_a, d_b, d_c, N);  
    cudaMemcpy(h_c, d_c, size, cudaMemcpyDeviceToHost);  
  
    cudaFree(d_a); cudaFree(d_b); cudaFree(d_c);  
    free(h_a); free(h_b);  
    return 0;}
```




GPU CODE WORKFLOW

Invoke the CUDA Kernel

```
int main(void) {  
    size_t size = N * sizeof(int);  
    int *h_a, *h_b; int *d_a, *d_b, *d_c;  
    h_a = (int *)malloc(size);  
    h_b = (int *)malloc(size);  
    ...  
    cudaMalloc((void **)&d_a, size);  
    cudaMalloc((void **)&d_b, size);  
    cudaMalloc((void **)&d_c, size);  
  
    cudaMemcpy(d_a, h_a, size, cudaMemcpyHostToDevice);  
    cudaMemcpy(d_b, h_b, size, cudaMemcpyHostToDevice);  
    vectorAdd<<<grid, block>>>>(d_a, d_b, d_c, N);  
    cudaMemcpy(h_c, d_c, size, cudaMemcpyDeviceToHost);  
  
    cudaFree(d_a); cudaFree(d_b); cudaFree(d_c);  
    free(h_a); free(h_b);  
    return 0;}
```


Device Pointers



GPU CODE WORKFLOW

Copy result from GPU to CPU

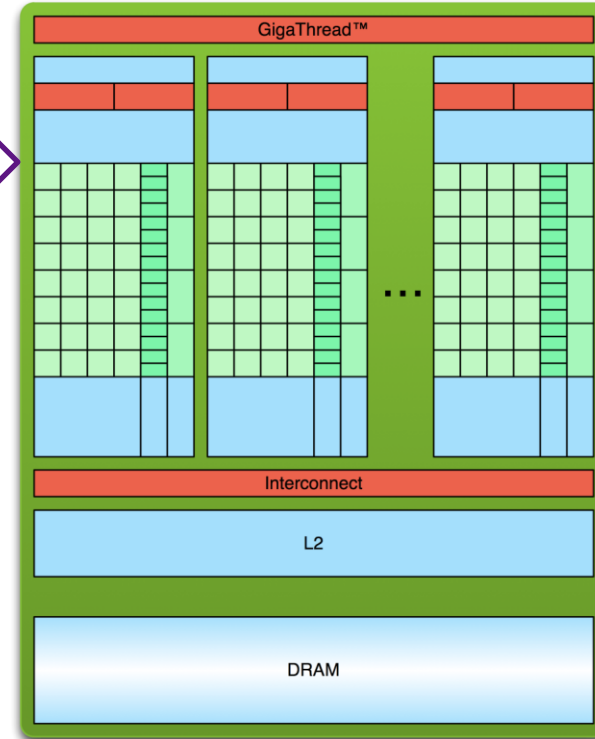

```
int main(void) {  
    size_t size = N * sizeof(int);  
    int *h_a, *h_b; int *d_a, *d_b, *d_c;  
    h_a = (int *)malloc(size);  
    h_b = (int *)malloc(size);  
    ...  
    cudaMalloc((void **)&d_a, size);  
    cudaMalloc((void **)&d_b, size);  
    cudaMalloc((void **)&d_c, size);  
  
    cudaMemcpy(d_a, h_a, size, cudaMemcpyHostToDevice);  
    cudaMemcpy(d_b, h_b, size, cudaMemcpyHostToDevice);  
    vectorAdd<<<grid, block>>>(d_a, d_b, d_c, N);  
    cudaMemcpy(h_c, d_c, size, cudaMemcpyDeviceToHost);  
  
    cudaFree(d_a); cudaFree(d_b); cudaFree(d_c);  
    free(h_a); free(h_b);  
    return 0;}
```




GPU CODE WORKFLOW

Release GPU Memories

```
int main(void) {  
    size_t size = N * sizeof(int);  
    int *h_a, *h_b; int *d_a, *d_b, *d_c;  
    h_a = (int *)malloc(size);  
    h_b = (int *)malloc(size);  
    ...  
    cudaMalloc((void **)&d_a, size);  
    cudaMalloc((void **)&d_b, size);  
    cudaMalloc((void **)&d_c, size);  
  
    cudaMemcpy(d_a, h_a, size, cudaMemcpyHostToDevice);  
    cudaMemcpy(d_b, h_b, size, cudaMemcpyHostToDevice);  
    vectorAdd<<<grid, block>>>(d_a, d_b, d_c, N);  
    cudaMemcpy(h_c, d_c, size, cudaMemcpyDeviceToHost);  
  
    cudaFree(d_a); cudaFree(d_b); cudaFree(d_c);  
    free(h_a); free(h_b);  
    return 0;}
```



图像处理




Pixel of an RGB image are formed from the corresponding pixel of the three component images

Sobel 边缘检测



$$\mathbf{G}_x = \begin{bmatrix} +1 & 0 & -1 \\ +2 & 0 & -2 \\ +1 & 0 & -1 \end{bmatrix} * \mathbf{A} \quad \text{and} \quad \mathbf{G}_y = \begin{bmatrix} +1 & +2 & +1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix} * \mathbf{A}$$

Sobel 边缘检测



更多资源：

<https://developer.nvidia-china.com>



何琨-Ken

北京 密云



扫一扫上面的二维码图案，加我微信

<https://www.nvidia.cn/developer/community-training/>

